



# BIG DATA ANALYTICS

Final Report

Apache Hbase

Lubaina Navaid  
10966

# Big Data Analytics

*Final Project: Report*

## About The Dataset

The dataset contains e-commerce data for the month of October 2019, from a large, multi-category online store. Each row within the dataset represents a purchase event, with a many-to-many relationship between products and users. This means that each row has to do with a specific product that was purchased as a specific point in time, with the time of purchase being used as the row key. Multiple purchase events can be carried out by the same user.

This dataset was taken from [Kaggle](#).

## Use Cases of Big Data in E-commerce

Big data tools can be used in e-commerce for a much deeper analysis than would be possible using traditional methods. Data analytics tools for big data are used for

- Optimizing back-office processes
- Enhancing customer experience
- Streamlining supply chain
- Identifying and preventing fraud
- Improved position in market against competitors due to optimization in pricing
- Demand forecasting

## HBase Commands

### Setting Up Hbase

#### Downloading from git

Git clone <https://github.com/big-data-europe/docker-hbase>

#### Using docker compose to set up container:

```
D:\big.data\docker-hbase>docker-compose -f docker-compose-distributed-local.yml up --d
```

#### Using docker ps to see running containers on network

```
D:\big.data\docker-hbase>docker ps
```

## Big data analytics

```
D:\big.data\docker-hbase>docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS
3473bcc186d9	bde2020/hadoop-datanode:2.0.0-hadoop2.7.4-java8	"/entrypoint.sh /run..."	16 hours ago	Up 51 seconds (healthy)	0.0.0.0:50075->50075/tcp
2aa106df9c71	bde2020/hadoop-namenode:2.0.0-hadoop2.7.4-java8	"/entrypoint.sh /run..."	16 hours ago	Up 50 seconds (healthy)	0.0.0.0:50070->50070/tcp
a843a9346a84	bde2020/hadoop-historyserver:2.0.0-hadoop2.7.4-java8	"/entrypoint.sh /run..."	16 hours ago	Up 52 seconds (health: starting)	0.0.0.0:8188->8188/tcp
807989bab4a1	bde2020/hadoop-nodemanager:2.0.0-hadoop2.7.4-java8	"/entrypoint.sh /run..."	16 hours ago	Up 54 seconds (health: starting)	0.0.0.0:8042->8042/tcp
b50a26ef862b	bde2020/hbase-regionserver:1.0.0-hbase1.2.6	"/entrypoint.sh /run..."	16 hours ago	Up 55 seconds	16020/tcp, 0.0.0.0:16030->16030/tcp
0cdb38263fdb	bde2020/hbase-master:1.0.0-hbase1.2.6	"/entrypoint.sh /run..."	16 hours ago	Up 48 seconds	16000/tcp, 0.0.0.0:16010->16010/tcp
a9d694d72275	bde2020/hadoop-resourcemanager:2.0.0-hadoop2.7.4-java8	"/entrypoint.sh /run..."	16 hours ago	Up 57 seconds (health: starting)	0.0.0.0:8088->8088/tcp
2ee686a891a9	zookeeper:3.4.10	"/docker-entrypoint..."	16 hours ago	Up 57 seconds	2888/tcp, 0.0.0.0:2181->2181/tcp, 3888/tcp

### Copying data into hbase master & namenode containers

```
D:\big.data\docker-hbase>docker cp 2019-Oct.csv 0cdb38263fdb:/hadoop-data
```

```
D:\big.data\docker-hbase>docker cp 2019-Oct.csv 2aa106df9c71:/hadoop-data
```

### executing hbase master:

```
D:\big.data\docker-hbase>docker exec -it 0cdb38263fdb /bin/bash
```

### Entering hbase shell:

```
root@hbase-master: hbase shell
```

### Viewing tables:

```
hbase(main):005:0> list
```

### Creating table:

```
hbase(main):006:0> create 'ecommerce-  
data','event_type','event','product','category','brand_price','user'
```

### Importing data from container to hbase

```
root@hbase-master: /# hbase org.apache.hadoop.hbase.mapreduce.ImportTsv -Dimporttsv.separator=''  
-
```

```
Dimporttsv.columns=HBASE_ROW_KEY,event:event_type,product:product_id,category:category_id,category:category_code,brand_price:brand,brand_price:price,user:user_id,user:user_session ecommerce-  
data /hadoop-data/2019-Oct.csv
```

## Basic Commands

### Check if table exists

```
hbase(main):012:0> exists 'ecommerce-data'
```

```
hbase(main):012:0> exists 'ecommerce-data'  
Table ecommerce-data does exist  
0 row(s) in 0.0120 seconds
```

### Check if table is enabled

```
hbase(main):001:0> is_enabled 'ecommerce-data'
```

```
hbase(main):001:0> is_enabled 'ecommerce-data'
true
0 row(s) in 0.9340 seconds
```

### Disable and re-enable table

```
hbase(main):004:0> disable 'ecommerce-data'
```

```
hbase(main):005:0> is_enabled 'ecommerce-data'
```

```
hbase(main):006:0> enable 'ecommerce-data'
```

```
hbase(main):007:0> is_enabled 'ecommerce-data'
```

```
hbase(main):004:0> disable 'ecommerce-data'
0 row(s) in 8.3480 seconds

hbase(main):005:0> is_enabled 'ecommerce-data'
false
0 row(s) in 0.0240 seconds

hbase(main):006:0> enable 'ecommerce-data'
0 row(s) in 5.3430 seconds

hbase(main):007:0> is_enabled 'ecommerce-data'
true
0 row(s) in 0.0270 seconds
```

### Describe table

```
hbase(main):002:0> describe 'ecommerce-data'
```

```
hbase(main):002:0> describe 'ecommerce-data'
Table ecommerce-data is ENABLED
ecommerce-data
COLUMN FAMILIES DESCRIPTION
{NAME => 'brand_price', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS => '0', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
{NAME => 'category', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS => '0', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
{NAME => 'event', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS => '0', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
```

### Counting values within table

```
hbase(main):005:0> Count 'ecommerce-data'
```

```
hbase(main):005:0> count 'ecommerce-data', CACHE=>1000
Current count: 1000, row: 2019-10-01 02:19:13 UTC
Current count: 2000, row: 2019-10-01 02:35:54 UTC
Current count: 3000, row: 2019-10-01 02:52:34 UTC
Current count: 4000, row: 2019-10-01 03:09:14 UTC
Current count: 5000, row: 2019-10-01 03:25:54 UTC
```

#### Altering table to add new column called 'column\_new'

```
hbase(main):003:0> alter 'ecommerce-data', NAME='column_new', VERSIONS=>5
```

```
hbase(main):003:0> alter 'ecommerce-data', NAME='column_new', VERSIONS=>5
Updating all regions with the new schema...
0/1 regions updated.
1/1 regions updated.
Done.
```

#### Scanning table to check if imported correctly

```
hbase(main):018:0> scan 'ecommerce-data',{FILTER=>"PageFilter(2)"}
```

```
hbase(main):018:0> scan 'ecommerce-data',{FILTER=>"PageFilter(2)"}
```

ROW	COLUMN+CELL
2019-10-01 00:00:00 U	column=brand_price:brand, timestamp=1653911448472, value=aqua TC
2019-10-01 00:00:00 U	column=brand_price:price, timestamp=1653911448472, value=33.2 TC
2019-10-01 00:00:00 U	column=category:category_code, timestamp=1653911448472, value=TC
2019-10-01 00:00:00 U	column=category:category_id, timestamp=1653911448472, value=2 TC

#### Getting data from row number '2019-01-01 00:00:08 UTC'

```
hbase(main):001:0> get 'ecommerce-data','2019-10-01 00:00:08 UTC'
```

```
hbase(main):001:0> get 'ecommerce-data','2019-10-01 00:00:08 UTC'
```

COLUMN	CELL
brand_price:brand	timestamp=1653911448472, value=luminarc
brand_price:price	timestamp=1653911448472, value=41.16
category:category_code	timestamp=1653911448472, value=
e	
category:category_id	timestamp=1653911448472, value=2.05301E+18
event:event_type	timestamp=1653911448472, value=view
product:product_id	timestamp=1653911448472, value=31500053
user:user_id	timestamp=1653911448472, value=550978835
user:user_session	timestamp=1653911448472, value=6280d577-25c8-4147-99a7-abc604

#### Getting data from column 'brand' in column family brand\_price for '2019-10-01 08:59:06 UTC'

```
hbase(main):002:0> get 'ecommerce-data', '2019-10-01 08:59:06 UTC',
{COLUMNS=>'brand_price:brand'}
```

```
hbase(main):002:0> get 'ecommerce-data', '2019-10-01 08:59:06 UTC', {COLUMNS=>'brand_price:brand'}
COLUMN          CELL
brand_price:brand timestamp=1653911448472, value=goodyear
```

#### Getting data from columns 'price' and 'category\_id' for '2019-10-01 07:11:42 UTC'

```
hbase(main):002:0> get 'ecommerce-data','2019-10-01 07:11:42 UTC',
{COLUMNS=>['brand_price:price','category:category_id']}
```

```
hbase(main):002:0> get 'ecommerce-data','2019-10-01 07:11:42 UTC', {COLUMNS=>['brand_price:price','category:category_id']}
COLUMN          CELL
brand_price:price timestamp=1653911448472, value=73.36
category:category_id timestamp=1653911448472, value=2.05301E+18
```

#### Getting timeline and values for row '2019-10-01 09:03:50 UTC'

```
hbase(main):003:0> get 'ecommerce-data', '2019-10-01 09:03:50 UTC',{CONSISTENCY=> 'TIMELINE'}
```

```
hbase(main):003:0> get 'ecommerce-data', '2019-10-01 09:03:50 UTC',{CONSISTENCY=> 'TIMELINE'}
COLUMN          CELL
brand_price:brand timestamp=1653911448472, value=artel
brand_price:price timestamp=1653911448472, value=139.69
category:category_cod timestamp=1653911448472, value=
e
category:category_id timestamp=1653911448472, value=2.05301E+18
event:event_type   timestamp=1653911448472, value=view
product:product_id timestamp=1653911448472, value=2601938
user:user_id        timestamp=1653911448472, value=554541876
user:user_session   timestamp=1653911448472, value=c905ecce-d7a9-4653-bb71-1756c6b9ff5a
```

#### Deleting row '2019-10-01 00:00:04 UTC' from table

```
hbase(main):002:0> delete 'ecommerce-data', '2019-10-01 00:00:04 UTC', 'event:event_type'
```

```
hbase(main):002:0> delete 'ecommerce-data', '2019-10-01 00:00:04 UTC', 'event:event_type'
0 row(s) in 1.1560 seconds
```

#### Get value in row '2019-10-01 00:00:00 UTC' where value 'shiseido' is included

```
hbase(main):005:0> get 'ecommerce-data','2019-10-01 00:00:00 UTC',{COLUMN=>'brand_price:brand',ATTRIBUTES=>{'2019-10-01 00:00:00 UTC'=>'shiseido'}}
```

```
hbase(main):005:0> get 'ecommerce-data','2019-10-01 00:00:00 UTC',{COLUMN=>'brand_price:brand',ATTRIBUTES=>{'2019-10-01 00:00:00 UTC'=>'shiseido'}}
COLUMN          CELL
brand_price:brand timestamp=1653911448472, value=aqua
1 row(s) in 0.3700 seconds
```

**Get counter cell value for row '2019-10-01 11:12:44 UTC', column 'category\_code' in column family 'category'**

```
hbase(main):006:0> get_counter 'ecommerce-data','2019-10-01 11:12:44 UTC','category:category_code'
```

```
hbase(main):006:0> get_counter 'ecommerce-data','2019-10-01 11:12:44 UTC','category:category_code'
COUNTER VALUE = 7308327773145296750
```

**See total number of splits in data on hbase by region**

```
hbase(main):007:0> get_splits 'ecommerce-data'
```

```
hbase(main):007:0> get_splits 'ecommerce-data'
Total number of splits = 1
```

**Put value '200' in price column for row '2019-10-01 16:07:31 UTC'**

```
hbase(main):001:0> put 'ecommerce-data','2019-10-01 16:07:31 UTC','brand_price:price','200'
```

```
hbase(main):001:0> put 'ecommerce-data','2019-10-01 16:07:31 UTC','brand_price:price','200'
0 row(s) in 1.4930 seconds
```

**Verifying put command using get command**

```
hbase(main):002:0> get 'ecommerce-data','2019-10-01 16:07:31 UTC','brand_price:price'
```

```
hbase(main):002:0> get 'ecommerce-data','2019-10-01 16:07:31 UTC','brand_price:price'
COLUMN                                CELL
brand_price:price                      timestamp=1654021644801, value=200
1 row(s) in 0.3430 seconds
```

**Perform meta scan on table**

```
hbase(main):003:0> scan 'hbase:meta'
```

```
scan 'hbase:meta'
ROW                                COLUMN+CELL
data,,1653895421351.d59a9d5ad39f35d8caafe279db1be490.  column=info:regioninfo, timestamp=1654020521935, value={ENCOD
59a9d5ad39f35d8caafe279db1be490.  ED => d59a9d5ad39f35d8caafe279db1be490, NAME => 'data,,165389
79db1be490.  5421351.d59a9d5ad39f35d8caafe279db1be490.', STARTKEY => '', EN
NDKEY => ''}
```

**Locating region of given row**

```
hbase(main):007:0> locate_region 'ecommerce-data','2019-10-01 11:12:44 UTC'
```

```
hbase(main):007:0> locate_region 'ecommerce-data', '2019-10-01 11:12:44 UTC'
HOST                REGION
hbase-region:16020  {ENCODED => bba77fe992adfd9f37cdf57003bfab02, NAME => 'ecommerce-data,,1653876390428.bba77fe992adfd9f37cdf57003bfab02.', STARTKEY => '', ENDKEY => ''}
1 row(s) in 0.0520 seconds
```

**Deleting cell value of column product\_id in row '2019-10-01 07:31:53 UTC' and verifying deleting using get command**

```
hbase(main):010:0> delete 'ecommerce-data', '2019-10-01 07:31:53 UTC', 'product:product_id'
```

```
hbase(main):011:0> get 'ecommerce-data', '2019-10-01 07:31:53 UTC'
```

```
hbase(main):010:0> delete 'ecommerce-data', '2019-10-01 07:31:53 UTC', 'product:product_id'
0 row(s) in 2.2670 seconds

hbase(main):011:0> get 'ecommerce-data', '2019-10-01 07:31:53 UTC'
COLUMN                CELL
brand_price:brand      timestamp=1653911448472, value=lenovo
brand_price:price      timestamp=1653911448472, value=1103.08
category:category_cod timestamp=1653911448472, value=computers.notebook
e
category:category_id   timestamp=1653911448472, value=2.05301E+18
event:event_type       timestamp=1653911448472, value=view
user:user_id           timestamp=1653911448472, value=512393839
user:user_session      timestamp=1653911448472, value=50d9c8ff-f367-40d7-a147-946a1634f38e
7 row(s) in 1.7530 seconds
```

**Get values for columns brand, user session and product\_id for row '2019-10-01 17:27:37 UTC'**

```
hbase(main):013:0> get 'ecommerce-data', '2019-10-01 17:27:37 UTC', {COLUMN=>['product:product_id', 'brand_price:brand', 'user:user_session']}
```

```
hbase(main):013:0> get 'ecommerce-data', '2019-10-01 17:27:37 UTC', {COLUMN=>['_product:product_id', 'brand_price:brand', 'user:user_session']}
COLUMN                CELL
brand_price:brand      timestamp=1653911448472, value=dymatize
product:product_id     timestamp=1653911448472, value=29501936
user:user_session      timestamp=1653911448472, value=dc688388-8570-4c09-add1-30e5407e3d92
3 row(s) in 0.0680 seconds
```

**Scan table for columns brand, price and user\_id starting from row '2019-10-01 00:01:15 UTC' with limit 10**

```
hbase(main):015:0> scan 'ecommerce-data', {COLUMNS=>['brand_price:brand', 'brand_price:price', 'user:user_id'], LIMIT=>10, STARTROW=>'2019-10-01 00:01:15 UTC'}
```



```
hbase(main):015:0> scan 'ecommerce-data',{COLUMNS=>['brand_price:brand','brand_price:price','user:user_id'], LIMIT=>10, STARTROW=>'2019-10-01 00:01:15 UTC'}
ROW COLUMN+CELL
2019-10-01 00:01:15 U column=brand_price:brand, timestamp=1653911448472, value=lg
TC
2019-10-01 00:01:15 U column=brand_price:price, timestamp=1653911448472, value=462.
TC 25
2019-10-01 00:01:15 U column=user:user_id, timestamp=1653911448472, value=537918940
TC
2019-10-01 00:01:16 U column=brand_price:brand, timestamp=1653911448472, value=noki
TC a
```

### Scan table for row '2019-10-01 14:37:02 UTC' using RowFilter

```
hbase(main):018:0> scan 'ecommerce-data',{FILTER=> "RowFilter(=,'binary:2019-10-01 14:37:02 UTC')"}

```

```
hbase(main):018:0> scan 'ecommerce-data',{FILTER=> "RowFilter(=,'binary:2019-10-01 14:37:02 UTC')"}
ROW COLUMN+CELL
2019-10-01 14:37:02 U column=brand_price:brand, timestamp=1653911448472, value=tim
TC
2019-10-01 14:37:02 U column=brand_price:price, timestamp=1653911448472, value=16.4
TC 5
2019-10-01 14:37:02 U column=category:category_code, timestamp=1653911448472, value
TC =
2019-10-01 14:37:02 U column=category:category_id, timestamp=1653911448472, value=2
TC .05301E+18
2019-10-01 14:37:02 U column=event:event_type, timestamp=1653911448472, value=view
```

### Scan table for rows where brand=aqua with limit of 10 using value filter

```
hbase(main):019:0> scan 'ecommerce-data',{COLUMNS=>'brand_price:brand', LIMIT=>10, FILTER=>"ValueFilter(=,'binary:aqua')"}

```

```
hbase(main):019:0> scan 'ecommerce-data',{COLUMNS=>'brand_price:brand', LIMIT=>10, FILTER=>"ValueFilter(=,'binary:aqua')"}
ROW COLUMN+CELL
2019-10-01 00:00:00 U column=brand_price:brand, timestamp=1653911448472, value=aqua
TC
2019-10-01 00:11:29 U column=brand_price:brand, timestamp=1653911448472, value=aqua
TC
2019-10-01 03:30:52 U column=brand_price:brand, timestamp=1653911448472, value=aqua
TC
2019-10-01 03:32:04 U column=brand_price:brand, timestamp=1653911448472, value=aqua
TC
```

### Scan table for rows where price is less than 200 starting from row '2019-10-01 16:08:41 UTC', using ValueFilter

```
hbase(main):020:0> scan 'ecommerce-data',{COLUMNS=>'brand_price:price', LIMIT=>10, FILTER=>"ValueFilter(<,'binary:200')", STARTROW=>'2019-10-01 16:08:41 UTC'}
```

```
hbase(main):020:0> scan 'ecommerce-data',{COLUMNS=>'brand_price:price', LIMIT=>10, FILTER=>"ValueFilter(<,'binary:200')", STARTROW=>'2019-10-01 16:08:41 UTC'}
```

ROW	COLUMN+CELL
2019-10-01 16:08:43 U	column=brand_price:price, timestamp=1653911448472, value=18.0
TC	2
2019-10-01 16:08:45 U	column=brand_price:price, timestamp=1653911448472, value=153.
TC	67
2019-10-01 16:08:46 U	column=brand_price:price, timestamp=1653911448472, value=197.
TC	4
2019-10-01 16:08:47 U	column=brand_price:price, timestamp=1653911448472, value=180.
TC	49
2019-10-01 16:08:48 U	column=brand_price:price, timestamp=1653911448472, value=183.
TC	02
2019-10-01 16:08:51 U	column=brand_price:price, timestamp=1653911448472, value=1284

**Scan table for first key using FirstKeyOnlyFilter limiting rows to 10**

```
hbase(main):025:0> scan 'ecommerce-data',{LIMIT=>10,FILTER=>FirstKeyOnlyFilter.new()}
```

```
hbase(main):025:0> scan 'ecommerce-data',{LIMIT=>10,FILTER=>FirstKeyOnlyFilter.new()
}
```

ROW	COLUMN+CELL
2019-10-01 00:00:00 U	column=brand_price:brand, timestamp=1653911448472, value=aqua
TC	
2019-10-01 00:00:01 U	column=brand_price:brand, timestamp=1653911448472, value=leno
TC	vo
2019-10-01 00:00:04 U	column=brand_price:brand, timestamp=1653911448472, value=appl
TC	e
2019-10-01 00:00:05 U	column=brand_price:brand, timestamp=1653911448472, value=puls
TC	er
2019-10-01 00:00:08 U	column=brand_price:brand, timestamp=1653911448472, value=lumi
TC	narc

**Import PrefixFilter and Bytes utility, then scan table using prefix filter for row '2019-10-01 09:39:03 UTC'**

```
hbase(main):026:0> import org.apache.hadoop.hbase.filter.PrefixFilter
```

```
hbase(main):027:0> import org.apache.hadoop.hbase.util.Bytes
```

```
hbase(main):028:0> scan 'ecommerce-data',{FILTER=>PrefixFilter.new(Bytes.toBytes('2019-10-01
09:39:03 UTC'))}
```

```

hbase(main):026:0> import org.apache.hadoop.hbase.filter.PrefixFilter
=> Java::OrgApacheHadoopHbaseFilter::PrefixFilter
hbase(main):027:0> import org.apache.hadoop.hbase.util.Bytes
hbase(main):028:0> scan 'ecommerce-data',{FILTER=>PrefixFilter.new(Bytes.toBytes('20
19-10-01 09:39:03 UTC'))}
ROW          COLUMN+CELL
2019-10-01 09:39:03 U column=brand_price:brand, timestamp=1653911448472, value=inde
TC          sit
2019-10-01 09:39:03 U column=brand_price:price, timestamp=1653911448472, value=270.
TC          02
2019-10-01 09:39:03 U column=category:category_code, timestamp=1653911448472, value
TC          =appliances.kitchen.refrigerators
2019-10-01 09:39:03 U column=category:category_id, timestamp=1653911448472, value=2
TC          .05301E+18
2019-10-01 09:39:03 U column=event:event_type, timestamp=1653911448472, value=view
TC
2019-10-01 09:39:03 U column=product:product_id, timestamp=1653911448472, value=270
TC          1646

```

**Import ColumnPrefixFilter and scan table to get values of columns that start with 'b', from row '2019-10-01 07:47:34 UTC' to row '2019-10-01 07:47:47 UTC'**

```
hbase(main):029:0> import org.apache.hadoop.hbase.filter.ColumnPrefixFilter
```

```
hbase(main):031:0> scan 'ecommerce-data',{FILTER=>ColumnPrefixFilter.new(Bytes.toBytes('b')),
STARTROW=>'2019-10-01 07:47:34 UTC', STOPROW=>'2019-10-01 07:47:47 UTC'}
```

```

hbase(main):031:0> scan 'ecommerce-data',{FILTER=>ColumnPrefixFilter.new(Bytes.toBytes('b')), STARTROW=>'2019-10-01 07:47:34 UTC', STOPROW=>'2019-10-01 07:47:47 UTC'}
ROW          COLUMN+CELL
2019-10-01 07:47:34 U column=brand_price:brand, timestamp=1653911448472, value=phil
TC          ips
2019-10-01 07:47:35 U column=brand_price:brand, timestamp=1653911448472, value=huaw
TC          ei
2019-10-01 07:47:36 U column=brand_price:brand, timestamp=1653911448472, value=elen
TC          berg
2019-10-01 07:47:37 U column=brand_price:brand, timestamp=1653911448472, value=xiao
TC          mi

```

**Scan table for first 2 columns using ColumnCountGetFilter from row '2019-10-01 11:43:27 UTC' to row '2019-10-01 11:47:04 UTC'**

```
hbase(main):032:0> scan 'ecommerce-data',{STARTROW=>'2019-10-01 11:43:27 UTC',
STOPROW=>'2019-10-01 11:47:04 UTC',FILTER=>"ColumnCountGetFilter(2)"}

```

```
hbase(main):032:0> scan 'ecommerce-data',{STARTROW=>'2019-10-01 11:43:27 UTC', STOPROW=>'2019-10-01 11:47:04 UTC',FILTER=>"ColumnCountGetFilter(2)"}
ROW COLUMN+CELL
2019-10-01 11:43:27 U column=brand_price:brand, timestamp=1653911448472, value=good
TC year
2019-10-01 11:43:27 U column=brand_price:price, timestamp=1653911448472, value=117.
TC 63
2019-10-01 11:43:28 U column=brand_price:brand, timestamp=1653911448472, value=tp-1
TC ink
2019-10-01 11:43:28 U column=brand_price:price, timestamp=1653911448472, value=11.2
TC
```

**Scan table for 5 rows using PageFilter starting from row '2019-10-01 14:12:27 UTC' to row '2019-10-01 14:12:39 UTC'**

```
hbase(main):033:0> scan 'ecommerce-data',{STARTROW=>'2019-10-01 14:12:27 UTC',
STOPROW=>'2019-10-01 14:12:39 UTC', FILTER=>"PageFilter(5)"}

```

```
hbase(main):033:0> scan 'ecommerce-data',{STARTROW=>'2019-10-01 14:12:27 UTC', STOPROW=>'2019-10-01 14:12:39 UTC', FILTER=>"PageFilter(5)"}
ROW COLUMN+CELL
2019-10-01 14:12:27 U column=brand_price:brand, timestamp=1653911448472, value=daus
TC cher
2019-10-01 14:12:27 U column=brand_price:price, timestamp=1653911448472, value=483.
TC 9
2019-10-01 14:12:27 U column=category:category_code, timestamp=1653911448472, value
TC =
2019-10-01 14:12:27 U column=category:category_id, timestamp=1653911448472, value=2
TC .05301E+18
2019-10-01 14:12:27 U column=event:event_type, timestamp=1653911448472, value=view
```

**Scan table for rows starting from '2019-10-01 00:00:00 UTC' with limit 10 using InclusiveStopFilter**

```
hbase(main):034:0> scan 'ecommerce-data',{STARTROW=>'2019-10-01 00:00:00 UTC',
LIMIT=>10,FILTER=>"InclusiveStopFilter('row15')"}

```

```
hbase(main):034:0> scan 'ecommerce-data',{STARTROW=>'2019-10-01 00:00:00 UTC', LIMIT=>10,FILTER=>"InclusiveStopFilter('row15')"}
ROW COLUMN+CELL
2019-10-01 00:00:00 U column=brand_price:brand, timestamp=1653911448472, value=aqua
TC
2019-10-01 00:00:00 U column=brand_price:price, timestamp=1653911448472, value=33.2
TC
2019-10-01 00:00:00 U column=category:category_code, timestamp=1653911448472, value
TC =appliances.environment.water_heater
2019-10-01 00:00:00 U column=category:category_id, timestamp=1653911448472, value=2
TC .05301E+18
```

**Scan table to find values from column family which includes substring 'br' using FamilyFilter**

```
hbase(main):037:0> scan 'ecommerce-data',{STARTROW=>'2019-10-01 07:04:39 UTC',
STOPROW=>'2019-10-01 07:04:50', FILTER=>"FamilyFilter(=,'substring:br')"}

```

```
hbase(main):037:0> scan 'ecommerce-data',{STARTROW=>'2019-10-01 07:04:39 UTC', STOPR
OW=>'2019-10-01 07:04:50', FILTER=>"FamilyFilter(=,'substring:br')"}
ROW COLUMN+CELL
2019-10-01 07:04:39 U column=brand_price:brand, timestamp=1653911448472, value=
TC
2019-10-01 07:04:39 U column=brand_price:price, timestamp=1653911448472, value=159.
TC 33
2019-10-01 07:04:40 U column=brand_price:brand, timestamp=1653911448472, value=maxx
TC is
2019-10-01 07:04:40 U column=brand_price:price, timestamp=1653911448472, value=71.8
TC 2
2019-10-01 07:04:41 U column=brand_price:brand, timestamp=1653911448472, value=haie
TC r
2019-10-01 07:04:41 U column=brand_price:price, timestamp=1653911448472, value=282.
```

**Scan for key values within column that contains the string 'pri' using QualifierFilter starting from row '2019-10-01 12:24:57 UTC'**

```
hbase(main):001:0> scan 'ecommerce-data',{STARTROW=>'2019-10-01 12:24:57 UTC',
STOPROW=>'2019-10-01 12:25:01 UTC', FILTER=>"QualifierFilter(=,'regexstring:pri')"}

```

```
hbase(main):001:0> scan 'ecommerce-data',{STARTROW=>'2019-10-01 12:24:57 UTC', STOP
ROW=>'2019-10-01 12:25:01 UTC', FILTER=>"QualifierFilter(=,'regexstring:pri')"}
ROW COLUMN+CELL
2019-10-01 12:24:57 U column=brand_price:price, timestamp=1653911448472, value=57.9
TC 2
2019-10-01 12:24:58 U column=brand_price:price, timestamp=1653911448472, value=514.
TC 75
2019-10-01 12:24:59 U column=brand_price:price, timestamp=1653911448472, value=46.3
```

**Scan table for key values within column that includes substring 'pri', and value less than 200 using FamilyFilter, ValueFilter and AND operator.**

```
hbase(main):002:0> scan 'ecommerce-data',{STARTROW=>'2019-10-01 07:23:21 UTC',
LIMIT=>10,FILTER=>"(FamilyFilter(=,'substring:pri')) AND (ValueFilter(<,'binary:200'))"}

```

```
hbase(main):002:0> scan 'ecommerce-data',{STARTROW=>'2019-10-01 07:23:21 UTC', LIM
IT=>10,FILTER=>"(FamilyFilter(=,'substring:pri')) AND (ValueFilter(<,'binary:200'))"}
ROW COLUMN+CELL
2019-10-01 07:23:21 U column=brand_price:brand, timestamp=1653911448472, value=
TC
2019-10-01 07:23:21 U column=brand_price:price, timestamp=1653911448472, value=10.2
TC 4
2019-10-01 07:23:24 U column=brand_price:price, timestamp=1653911448472, value=179.
TC 35
2019-10-01 07:23:26 U column=brand_price:brand, timestamp=1653911448472, value=
TC
```



**Scan table for key values where column includes substring 'br' and value 'aq', OR where column includes substring 'pri' and value greater than 1000.**

```
hbase(main):005:0> scan 'ecommerce-data',{STARTROW=>'2019-10-01 04:12:13 UTC', LIMIT=>10,
FILTER=>"(FamilyFilter(=,'substring:br')) AND (ValueFilter(=,'binary:aqu')) OR
(FamilyFilter(=,'substring:pri')) AND (ValueFilter(>,'binary:1000'))"}
```

```
hbase(main):005:0> scan 'ecommerce-data',{STARTROW=>'2019-10-01 04:12:13 UTC', LIMIT
=>10, FILTER=>"(FamilyFilter(=,'substring:br')) AND (ValueFilter(=,'binary:aqu')) OR
(FamilyFilter(=,'substring:pri')) AND (ValueFilter(>,'binary:1000'))"}
```

ROW	COLUMN+CELL
2019-10-01 04:12:13 UTC	column=brand_price:brand, timestamp=1653911448472, value=apple
2019-10-01 04:12:13 UTC	column=brand_price:price, timestamp=1653911448472, value=2072.1
2019-10-01 04:12:14 UTC	column=brand_price:price, timestamp=1653911448472, value=87
2019-10-01 04:12:15 UTC	column=brand_price:brand, timestamp=1653911448472, value=tomahawk
2019-10-01 04:12:15 UTC	column=brand_price:price, timestamp=1653911448472, value=97.8

**Scan table for rows using CompareFilter, RowFilter and SubstringComparator to find rowkeys which include the substring '2019-10-01 07:' with a limit of 20 rows**

```
hbase(main):012:0> scan 'ecommerce-data',{FILTER=>
RowFilter.new(CompareFilter::CompareOp.valueOf('EQUAL'),SubstringComparator.new('2019-10-01
07:')), LIMIT=>20}
```

```
hbase(main):012:0> scan 'ecommerce-data',{FILTER=> RowFilter.new(CompareFilter::Comp
areOp.valueOf('EQUAL'),SubstringComparator.new('2019-10-01 07:')), LIMIT=>20}
```

ROW	COLUMN+CELL
2019-10-01 07:00:00 UTC	column=brand_price:brand, timestamp=1653911448472, value=msi
2019-10-01 07:00:00 UTC	column=brand_price:price, timestamp=1653911448472, value=463.31
2019-10-01 07:00:00 UTC	column=category:category_code, timestamp=1653911448472, value=computers.components.videocards
2019-10-01 07:00:00 UTC	column=category:category_id, timestamp=1653911448472, value=2053015110

**Put value 'sony' into brand column for row '2019-10-01 00:00:00 UTC'. Verifying put command using get command**

```
hbase(main):017:0> put 'ecommerce-data','2019-10-01 00:00:00 UTC', 'brand_price:brand','sony'
```

```
hbase(main):018:0> get 'ecommerce-data','2019-10-01 00:00:00 UTC'
```

```
hbase(main):017:0> put 'ecommerce-data','2019-10-01 00:00:00 UTC', 'brand_price:brand','sony'
0 row(s) in 0.3300 seconds

hbase(main):018:0> get 'ecommerce-data','2019-10-01 00:00:00 UTC'
COLUMN                                CELL
brand_price:brand                      timestamp=1654112317700, value=sony
brand_price:price                      timestamp=1653911448472, value=33.2
```

### Disable and drop table 'data'

```
hbase(main):008:0> list
```

```
hbase(main):009:0> disable 'data'
```

```
hbase(main):010:0> drop 'data'
```

```
hbase(main):008:0> list
TABLE
data
data2
ecommerce-data
3 row(s) in 0.2820 seconds

=> ["data", "data2", "ecommerce-data"]
hbase(main):009:0> disable 'data'
0 row(s) in 5.6700 seconds

hbase(main):010:0> drop 'data'
0 row(s) in 6.3770 seconds

hbase(main):011:0> list
TABLE
data2
ecommerce-data
2 row(s) in 0.0100 seconds
```

### Snapshotting table and checking list

```
hbase(main):017:0> snapshot 'ecommerce-data','snapshot1'
```

hbase(main):018:0> list\_snapshots

```
hbase(main):017:0> snapshot 'ecommerce-data', 'snapshot1'
0 row(s) in 6.9880 seconds

hbase(main):018:0> list_snapshots
SNAPSHOT          TABLE + CREATION TIME
snapshot1         ecommerce-data (Wed Jun 01 21:40:06 +0000 2022)
1 row(s) in 2.1650 seconds
```

### Using snapshot to clone table

hbase(main):020:0> clone\_snapshot 'snapshot1', 'ecommerce-data2'

hbase(main):021:0> list

```
hbase(main):020:0> clone_snapshot 'snapshot1', 'ecommerce-data2'
0 row(s) in 9.8920 seconds

hbase(main):021:0> list
TABLE
data2
ecommerce-data
ecommerce-data2
3 row(s) in 0.0120 seconds
```



### Attempted peers and replication

```
hbase(main):015:0> list_peers
  PEER_ID CLUSTER_KEY STATE TABLE_CFS
0 row(s) in 0.1290 seconds

hbase(main):016:0> add_peer '1', 'zoo:2181'
NoMethodError: undefined method `peer' for #<Object:0x62066ae7>

hbase(main):017:0> add_peer '1','zoo:2181'
0 row(s) in 1.9970 seconds

hbase(main):018:0> list_peers
NameError: undefined local variable or method `peers' for #<Object:0x62066ae7>

hbase(main):019:0> list_peers
  PEER_ID CLUSTER_KEY STATE TABLE_CFS
  1 zoo:2181 ENABLED
1 row(s) in 0.5600 seconds

hbase(main):020:0> get_peer_config
NameError: undefined local variable or method `get_peer_config' for #<Object:0x62066ae7>

hbase(main):021:0> enable_table_replication 'ecommerce-data'
2022-05-31 18:52:25,133 ERROR [main] replication.ReplicationPeersZKImpl: Can't get p
```