



WASILA REHMAN

12348

PROJECT 1 – KAFKA SPARK STREAMING PIPELINE

WASILA REHMAN - 12348

[Email address]

Industry

Health Care

Problem Statement

According to the World Health Organization, every year 12 million deaths occur worldwide due to heart diseases. Heart disease is one of the major causes of mortality among the population of the world.

The major challenge in heart disease is its detection. An early diagnosis of heart disease can help us make lifestyle changes in high-risk patients and ultimately reduce the chances of future complications.

There are instruments available which can predict heart disease but either are expensive or not fast enough to calculate chance of heart disease in humans.

That's why my goal in this project is to build an analytical dashboard on this big data so doctors can easily access insightful reports on key metrics.

Dataset Information

Originally, the [dataset](#) come from the CDC and is a major part of the Behavioral Risk Factor Surveillance System (BRFSS), which conducts annual telephone surveys to gather data on the health status of U.S. residents.

This dataset is already cleaned and has 319795 rows × 18 columns.

Following are the features present in this dataset:

1. **HeartDisease:** Respondents that have ever reported having coronary heart disease (CHD) or myocardial infarction (MI).
2. **BMI:** Body Mass Index (BMI).
3. **Smoking:** Have you smoked at least 100 cigarettes in your entire life?
4. **AlcoholDrinking:** Heavy drinkers (adult men having more than 14 drinks per week and adult women having more than 7 drinks per week)
5. **Stroke:** (Ever told) (you had) a stroke?
6. **PhysicalHealth:** Now thinking about your physical health, which includes physical illness and injury, for how many days during the past 30 days was your physical health not good? (0-30 days).
7. **MentalHealth:** Thinking about your mental health, for how many days during the past 30 days was your mental health not good? (0-30 days).
8. **DiffWalking:** Do you have serious difficulty walking or climbing stairs?
9. **Sex:** Are you male or female?
10. **AgeCategory:** Fourteen-level age category.
11. **Race:** Imputed race/ethnicity value.
12. **Diabetic:** (Ever told) (you had) diabetes?
13. **PhysicalActivity:** Adults who reported doing physical activity or exercise during the past 30 days other than their regular job.
14. **GenHealth:** Would you say that in general your health is...
15. **SleepTime:** On average, how many hours of sleep do you get in a 24-hour period?

16. **Asthma:** (Ever told) (you had) asthma?
17. **KidneyDisease:** Not including kidney stones, bladder infection or incontinence, were you ever told you had kidney disease?
18. **SkinCancer:** (Ever told) (you had) skin cancer?

Key Metrics to Track

According to BRFSS Machine Learning researchers, it was decided the key factors that contribute the most in heart diseases are:

- Mental Health
- Physical Health
- Age
- Gender
- BMI
- Sleep Time

For most hospitals the number of available doctors is always limited. Hence a quick diagnosis is needed to figure out whether a patient needs urgent treatment.

Analyzing dataset of the past 2 million records helps in uncovering the correlations between each health factor. Consequently, it saves time for doctor's diagnosis and understanding in what category a new patient will fall in.

The key questions this project answers are:

1. **Are the chances of Heart Disease higher in a certain Age category?**
2. **What is the average BMI of Heart Patients in each Age Category?**
3. **Is there a correlation between Heart Diseases and Gender?**
4. **What is the average number of sleeping hours in heart patients of each age category?**

STAGE 1 – Docker Compose

Building the Docker Compose File with the following services:

- Hive-server
- Hive-metastore
- Hive-metastore-postgresql
- Kafka
- Prometheus
- Grafana (dependency on Prometheus)
- Zeppelin (dependency on hive)

STAGE 2 – Dependencies

Downloading dependencies in Zeppelin:

To install kafka-python in Zeppelin shell

%sh

```
pip install kafka-python
```

These jar files were the essence of code running inside Zeppelin. Figuring each dependency out from maven is crucial to run the pipeline.

```
%sh
```

```
mkdir /zeppelin/dep
```

```
cd /zeppelin/dep && wget https://repo1.maven.org/maven2/org/apache/spark/spark-streaming-kafka-0-8-assembly_2.11/2.0.2/spark-streaming-kafka-0-8-assembly_2.11-2.0.2.jar
```

```
cd /zeppelin/dep && wget https://repo1.maven.org/maven2/org/apache/spark/spark-streaming-kafka-0-10_2.11/2.0.2/spark-streaming-kafka-0-10_2.11-2.0.2.jar
```

```
cd /zeppelin/dep && wget https://repo1.maven.org/maven2/org/apache/hive/hive-jdbc/2.3.6/hive-jdbc-2.3.6.jar
```

```
cd /zeppelin/dep && wget https://repo1.maven.org/maven2/org/apache/hive/hive-service/2.3.6/hive-service-2.3.6.jar
```

```
cd /zeppelin/dep && wget https://repo1.maven.org/maven2/org/apache/hive/hive-exec/2.3.6/hive-exec-2.3.6.jar
```

Downloading the same dependencies for spark, consumer.pyspark, and producer.pyspark

```
%spark.dep
```

```
z.reset()
```

```
z.load("/zeppelin/dep/hive-jdbc-2.3.6.jar")
```

```
z.load("/zeppelin/dep/hive-service-2.3.6.jar")
```

```
z.load("/zeppelin/dep/hive-exec-2.3.6.jar")
```

```
%producer.dep
```

```
z.reset()
```

```
z.load("/zeppelin/dep/spark-streaming-kafka-0-10_2.11-2.0.2.jar")
```

```
z.load("/zeppelin/dep/spark-sql-kafka-0-10_2.11:2.3.0.jar")
```

```
%consumer.dep
```

```
z.reset()
```

```
z.load("/zeppelin/dep/spark-sql-kafka-0-10_2.11:2.2.0.jar")
```

```
z.load("/zeppelin/dep/spark-streaming-kafka-0-10-assembly_2.11-2.0.2.jar")
z.load("/zeppelin/dep/kafka_2.11.jar")
z.load("/zeppelin/dep/kafka-clients.jar")
z.load("/zeppelin/dep/hive-jdbc-2.3.6.jar")
z.load("/zeppelin/dep/hive-service-2.3.6.jar")
z.load("/zeppelin/dep/hive-exec-2.3.6.jar")
```

STAGE 3 – Building Kafka Producer

(CAN KEEP RUNNING FOR 40 MINUTES BEFORE KAFKA PRODUCER TIMEOUTS)

First step is to build the data frame for the Producer by reading a big data csv file.

```
%producer.pyspark
```

```
df = (spark.read.format("com.databricks.spark.csv")
      .option("header", "true")
      .option("inferSchema", "true")
      .load("/datadrive/heart_2020.csv"))
```

```
df_list = df.collect()
df.show()
df.printSchema()
```

Next step is to enable Producer send data to our exposed Kafka_brokers Ports. Kafka_Brokers are the mid-level broker who take data from the Producer and send it to Consumer. It is responsible for all the connections.

In the Kafka Topic we have specified our key and string. The json dumps it into our row_dictionary.

```
%producer.pyspark

import time

import json

import random
```

```
import logging

from kafka import KafkaProducer
from kafka.errors import KafkaError

KAFKA_BROKER = "172.25.0.12:9092"
KAFKA_INPUT_TOPIC = "default_topic"

producer = KafkaProducer(bootstrap_servers=[KAFKA_BROKER])
index = 0

while True:

    row_dict = df_list[index].asDict()

    future = producer.send(
        topic=KAFKA_INPUT_TOPIC,
        key=str(row_dict["PatientID"]).encode("utf-8"),
        value=json.dumps(row_dict).encode("utf-8"))

    try:

        record_metadata = future.get(timeout=10)
    except KafkaError:
        # Decide what to do if produce request failed...
        logging.exception("Error")
        pass

    producer.flush()
```

```
index += 1
```

```
time.sleep(random.uniform(0.1,3.0))
```

```
return (row_dict)
```

Stage 4 – Building Kafka Consumer

TIME TAKEN: 8MINUTES (until zeppelin range ends)

Kafka Consumer builds a spark session and creates a table in Hive while consumer keeps on appending on each iteration.



```
%consumer.pyspark
```

```
import json
```

```
from pyspark.sql import SparkSession
```

```
from pyspark.sql import HiveContext
```

```
from pyspark.sql import Row
```

```
from kafka import KafkaConsumer
```

```
import pyspark
```

```
from ast import literal_eval
```

```
spark =
```

```
SparkSession.builder.master("local[1]").appName("Heart_Diseases").enableHiveSupport().getOrCreate()
```

```
HiveContext = HiveContext(sc)
```

```
HiveContext.sql("CREATE DATABASE IF NOT EXISTS heart_db")
```

```
# To consume latest messages and auto-commit offsets
```

```
consumer = KafkaConsumer('default_topic',  
                           group_id='final-consumer',  
                           bootstrap_servers=['172.25.0.12:9092'])
```

```
for message in consumer:
```

```
    data = literal_eval(message.value)  
    studentDf = spark.createDataFrame([  
        Row(PatientID= data['PatientID'], heartDisease=data['HeartDisease'],  
PhysicalHealth=data['PhysicalHealth'], MentalHealth=data['MentalHealth'],  
AgeCategory=data['AgeCategory'], SleepTime=data['SleepTime'], gender=data['Sex'], BMI=data['BMI'])  
    ])
```

```
    studentDf.write.mode('append').saveAsTable("heart_db.table_1")
```

```
    df1 = spark.sql("select * from heart_db.table_1")
```

```
    df1.show()
```

```
# consume earliest available messages, don't commit offsets
```

```
KafkaConsumer(auto_offset_reset='earliest', enable_auto_commit=False)
```

```
# consume json messages
```

```
KafkaConsumer(value_deserializer=lambda m: json.loads(m.decode('ascii')))
```

```
# StopIteration if no message after 1sec
```

```
KafkaConsumer(consumer_timeout_ms=100)
```


STAGE 5 – Accessing tables from Hive & Hive Query Language (Spark SQL)

First we will access the data from the Hive table.

(TIME TAKEN : 2 SECONDS)

```
%spark
```

```
import org.apache.spark.sql.hive.HiveContext
```

```
val HiveContext = new org.apache.spark.sql.hive.HiveContext(sc)
```

```
HiveContext.sql("CREATE DATABASE IF NOT EXISTS heart_patients")
```

```
HiveContext.sql("CREATE TABLE IF NOT EXISTS heart_patients_Dets(PatientID INT, HeartDisease STRING, MentalHealth STRING, PhysicalHealth STRING, BMI INT, Sex string, AgeCategory string, SleepTime INT) ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' LINES TERMINATED BY '\n' ")
```

```
HiveContext.sql("LOAD DATA LOCAL INPATH '/datadrive/table_1.csv' INTO TABLE heart_patients_Dets")
```

```
val df1 = HiveContext.sql("Select * from heart_patients_Dets").show()
```

OUTPUT:

warning: there was one deprecation warning; re-run with -deprecation for details

PatientID	HeartDisease	MentalHealth	PhysicalHealth	BMI	Sex	AgeCategory	SleepTime
1	No	30	3	16	Female	55-59	5
2	No	0	0	20	Female	80 or older	7
3	No	30	20	26	Male	65-69	8
4	No	0	0	24	Female	75-79	6
5	No	0	28	23	Female	40-44	8
6	Yes	0	6	28	Female	75-79	12
7	No	0	15	21	Female	70-74	4
8	No	0	5	31	Female	80 or older	9
9	No	0	0	26	Female	80 or older	5
10	No	0	0	40	Male	65-69	10
11	Yes	0	30	34	Male	60-64	15
12	No	0	0	28	Female	55-59	5
13	No	0	0	28	Male	75-79	8

Took 2 sec. Last updated by anonymous at June 04 2022, 12:57:00 AM. (outdated)

Next, we start querying the data to figure out dataframes for our Key Metrics.

Query 1 – to return patients with SleepTime < 10

(TIME TAKEN 5 SECS)

%spark

```
val df2 = HiveContext.sql("SELECT PatientID, HeartDisease FROM heart_patients_Details WHERE SleepTime < 10 ORDER BY PatientID").show()
```

	2	No
	2	No
	2	No
	2	No
	2	No
	2	No
	3	No
	3	No
	3	No
	3	No
	3	No
	3	No
	4	No
	4	No

+-----+-----+

only showing top 20 rows

df2: Unit = ()

Took 5 sec. Last updated by anonymous at June 04 2022, 12:57:16 AM.

Query 2 – To gauge what is an average sleep time of HeartDisease Patients

(TIME TAKEN 8 SECS)

```
val df3 = HiveContext.sql("SELECT AgeCategory, AVG(SleepTime), count(HeartDisease) FROM heart_patients_Details WHERE HeartDisease = 'Yes' AND SleepTime < 8 GROUP BY AgeCategory Order BY count(HeartDisease) desc").show()
```

Took 8 sec. Last updated by anonymous at June 03 2022, 10:31:46 PM.

(TIME TAKEN 4 SECONDS)

count(HeartDisease)	AgeCategory
16347	80 or older
14541	70-74
12303	65-69
12147	75-79
9981	60-64
6606	55-59
4149	50-54
2232	45-49
1458	40-44
888	35-39
678	30-34
399	25-29
390	18-24

Took 4 sec. Last updated by anonymous at June 03 2022, 8:04:27 PM. (outdated)

Query 4 – To see average BMI of Heart Patients based on AgeCategory

(TIME TAKEN 5 SECONDS)

```
val df5 = HiveContext.sql("Select AVG(BMI) , count(HeartDisease), AgeCategory from heart_patients_Details where HeartDisease = 'Yes' group by AgeCategory Order by AVG(BMI), count(HeartDisease) desc").show()
```

avg(BMI)	count(HeartDisease)	AgeCategory
25.53846153846154	390	18-24
26.406496604881628	16347	80 or older
27.586466165413533	399	25-29
28.125710051864658	12147	75-79
28.901382298328862	14541	70-74
29.623893805309734	678	30-34
29.708607656669106	12303	65-69
30.05743243243243	888	35-39
30.13195070634205	9981	60-64
30.409173478655767	6606	55-59
30.96529284164859	4149	50-54
31.409465020576132	1458	40-44
31.928763440860216	2232	45-49

-- --

Took 5 sec. Last updated by anonymous at June 03 2022, 8:14:50 PM.

Query 5 – To gauge Heart Patients based on gender

(TIME TAKEN 6 SECONDS)

```
val df6 = HiveContext.sql("Select Sex, count(HeartDisease) from heart_patients_Details where HeartDisease = 'Yes' group by Sex").show()
```

Sex	count(HeartDisease)
Female	44936
Male	64556

df6: Unit = ()

Took 6 sec. Last updated by anonymous at June 03 2022, 10:24:59 PM.

Analysis & Conclusion

Let's answer the questions we posed in the beginning of the report.

1. Are the chances of Heart Disease higher in a certain Age category?

As can be seen from this output, Age Categories over 55+ have higher number of heart disease patients it means AgeCategory has an impact on number of heart disease patients.

+-----+-----+	
count(HeartDisease) AgeCategory	
+-----+-----+	
	16347 80 or older
	14541 70-74
	12303 65-69
	12147 75-79
	9981 60-64
	6606 55-59
	4149 50-54
	2232 45-49
	1458 40-44
	888 35-39
	678 30-34
	399 25-29
	390 18-24
+-----+-----+	

Took 4 sec. Last updated by anonymous at June 03 2022, 8:04:27 PM. (outdated)

2. What is the average BMI of Heart Patients in each Age Category?

As can be seen from this output, the highest BMI is 32 for a heart disease patient who is 45-49 years old.

avg(BMI)	count(HeartDisease)	AgeCategory
25.53846153846154	390	18-24
26.406496604881628	16347	80 or older
27.586466165413533	399	25-29
28.125710051864658	12147	75-79
28.901382298328862	14541	70-74
29.623893805309734	678	30-34
29.708607656669106	12303	65-69
30.05743243243243	888	35-39
30.13195070634205	9981	60-64
30.409173478655767	6606	55-59
30.96529284164859	4149	50-54
31.409465020576132	1458	40-44
31.928763440860216	2232	45-49

Took 5 sec. Last updated by anonymous at June 03 2022, 8:14:50 PM.

3. Is there a correlation between Heart Diseases and Gender?

There is no such huge correlation as the % of which males and females present in the dataset is almost same.

Sex	count(HeartDisease)
Female	44936
Male	64556

df6: Unit = ()

Took 6 sec. Last updated by anonymous at June 03 2022, 10:24:59 PM.

4. What is the average number of sleeping hours less than 8 in heart patients of each age category?

As the output displays, for heart patients who sleep less than 8 hours count of heart diseases are higher.

AgeCategory	avg(SleepTime)	count(HeartDisease)
70-74	6.159740754860846	10492
80 or older	6.165111561866126	9860
65-69	6.071546052631579	9728
60-64	5.922201138519924	8432
75-79	6.184158415841584	8080
55-59	5.848262032085562	5984
50-54	5.665955176093917	3748
45-49	5.622137404580153	2096
40-44	5.508241758241758	1456
35-39	5.4739336492890995	844
30-34	5.714285714285714	644
25-29	5.590909090909091	352
18-24	5.764705882352941	272

Took 8 sec. Last updated by anonymous at June 03 2022, 10:31:46 PM.

To conclude, all of these exploratory analysis can help doctors in providing a diagnosis within a few minutes which will save cost and prevent death cases through Heart Disease risk.